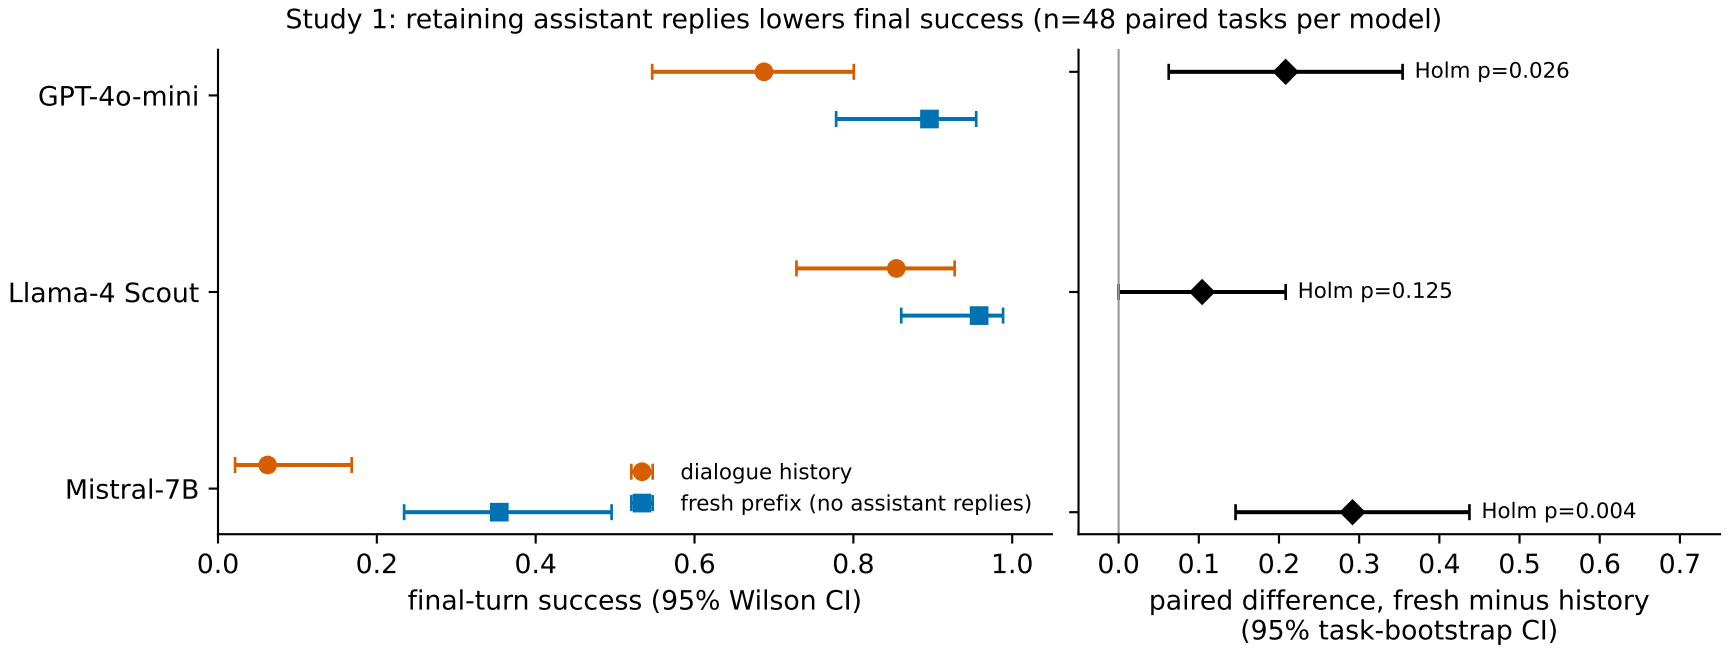
Study 1: retaining assistant replies lowers final success (n=48 paired tasks per model)
GPT-4o-mini
Llama-4 Scout
Mistral-7B
dialogue history
fresh prefix (no assistant replies)
0.0
0.2
0.4
0.6
0.8
1.0
final-turn success (95% Wilson CI)
Holm p=0.026
Holm p=0.125
Holm p=0.004
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
paired difference, fresh minus history
(95% task-bootstrap CI)